

Figure 4.4: Inference Latency Comparison in Milliseconds per Batch

